

# Assignment 3 - EN3160

Name Pathirana R. P. S.  
Index Number 210451U  
Github Link [https://github.com/RPX2001/Assignment3\\_EN3160.git](https://github.com/RPX2001/Assignment3_EN3160.git)

## 1 Changes for a single dense layer

### 1.1 Add middle layer with 100 nodes and sigmoid activation

In the original network architecture, there was only a single linear layer that directly mapped the input to the output classes. To improve the network's ability to learn complex patterns and non-linear relationships in the data, we added a hidden (middle) layer with 100 nodes and used the sigmoid function as the activation.

The sigmoid activation function was chosen for the middle layer because it introduces non-linearity, which enables the network to model more complex data patterns. The hidden layer transforms the input features into an intermediate representation, which then gets passed to the output layer for classification. This intermediate representation helps the network to better capture the patterns in the CIFAR-10 dataset, which includes various object categories (e.g., cars, animals, and ships).

Adding a hidden layer is a significant change as it transforms the model from a simple linear classifier to a two-layer neural network, which can better capture relationships within the data. This typically results in improved accuracy, albeit with slightly more computational cost and training time due to the increased number of parameters.



Figure 1: Single Dense Layer Output



Figure 2: After Adding middle layer with 100 nodes and sigmoid activation

```
1 Din = 3 * 32 * 32      # Input size (flattened CIFAR-10 image size)
2 Dhidden = 100          # Number of nodes in the hidden layer
```

```

3 K = 10 # Output size (number of classes in CIFAR-10)
4 std = 1e-5
5
6 w1 = torch.randn(Din, Dhidden) * std # Weights from input to hidden layer
7 b1 = torch.zeros(Dhidden) # Biases for the hidden layer
8 w2 = torch.randn(Dhidden, K) * std # Weights from hidden to output layer
9 b2 = torch.zeros(K)
10
11 h1 = torch.sigmoid(x_train.mm(w1) + b1)
12 y_pred = h1.mm(w2) + b2 # Hidden to output layer
13
14 dy_pred = (torch.softmax(y_pred, dim=1) - nn.functional.one_hot(labels, K).float()) /
15 Ntr
16 dw2 = h1.t().mm(dy_pred) + reg * w2
17 db2 = dy_pred.sum(dim=0)
18
19 # Gradients for hidden layer
20 dh1 = dy_pred.mm(w2.t()) * h1 * (1 - h1) # Derivative of sigmoid
21 dw1 = x_train.t().mm(dh1) + reg * w1
22 db1 = dh1.sum(dim=0)
23 w1 -= lr * dw1
24 b1 -= lr * db1
25 w2 -= lr * dw2
26 b2 -= lr * db2

```

## 1.2 Use cross entropy loss

In the original code, Mean Squared Error (MSE) was used as the loss function. However, for classification tasks, especially with multiple classes, Cross-Entropy Loss is generally preferred. Cross-Entropy Loss is better suited for classification because it penalizes incorrect class predictions more effectively, helping the model learn to make more confident predictions in the correct class. It does this by comparing the predicted probabilities with the actual class labels in a way that encourages the model to increase the probability for the true class while decreasing it for other classes.

Using Cross-Entropy Loss in this neural network setup should improve both convergence speed and accuracy, as the loss function will be more directly aligned with the objective of maximizing classification accuracy.

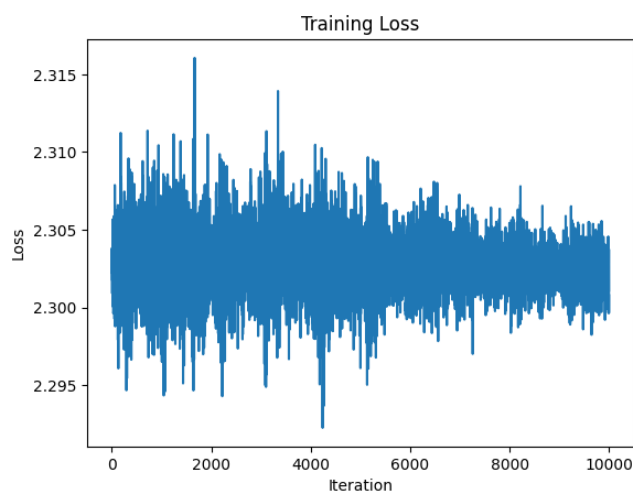


Figure 3: After using cross entropy loss

```

1 criterion = nn.CrossEntropyLoss()
2 h1 = torch.sigmoid(x_train.mm(w1) + b1)
3 y_pred = h1.mm(w2) + b2
4
5 # Calculate Cross-Entropy Loss
6 loss = criterion(y_pred, labels)

```

```

7 loss_history.append(loss.item())
8 running_loss += loss.item()
9 loss.backward()
10
11 # Update weights and biases manually (if not using an optimizer)
12 w1 -= lr * w1.grad
13 b1 -= lr * b1.grad
14 w2 -= lr * w2.grad
15 b2 -= lr * b2.grad

```

Results:

```

1 Epoch 1/10, Loss: 2.3029, Training Accuracy: 9.63%
2 Epoch 2/10, Loss: 2.3029, Training Accuracy: 9.91%
3 Epoch 3/10, Loss: 2.3028, Training Accuracy: 9.90%
4 Epoch 4/10, Loss: 2.3028, Training Accuracy: 9.93%
5 Epoch 5/10, Loss: 2.3028, Training Accuracy: 10.00%
6 Epoch 6/10, Loss: 2.3028, Training Accuracy: 10.07%
7 Epoch 7/10, Loss: 2.3027, Training Accuracy: 9.86%
8 Epoch 8/10, Loss: 2.3027, Training Accuracy: 9.89%
9 Epoch 9/10, Loss: 2.3025, Training Accuracy: 10.02%
10 Epoch 10/10, Loss: 2.3023, Training Accuracy: 10.82%
11
12 Test accuracy: 10.22%

```

## 2 LeNet-5 network for MNIST using Pytorch

Implemented a Convolutional Neural Network (CNN) inspired by the classic LeNet-5 architecture using PyTorch to classify images in the MNIST dataset. The network is designed with two convolutional layers followed by two fully connected layers to progressively extract features from the input images and classify them. The first convolutional layer detects low-level features with 6 filters, followed by a max-pooling layer to reduce spatial dimensions, while the second convolutional layer extracts more complex features with 16 filters. These features are then flattened and passed through two fully connected layers, which use a tanh activation to create a non-linear mapping from the extracted features to the ten-digit classes (0-9) of MNIST.

After defining the model architecture, we proceeded with training it for 10 epochs using the cross-entropy loss function and the Adam optimizer. This combination was chosen to efficiently optimize the model and minimize classification error. During each epoch, the model learned to adjust its parameters based on the calculated gradients, helping it distinguish between different digit classes with increasing accuracy. By the end of the training, we evaluated the model on both the training and test datasets to measure its generalization performance.

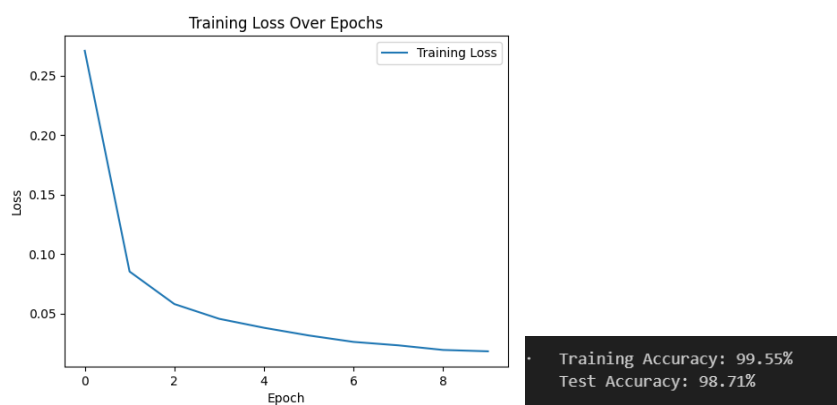


Figure 4: result of the model

```

1 import torch
2 import torch.nn as nn
3 import torch.optim as optim
4 import torchvision
5 import torchvision.transforms as transforms

```

```

6 import matplotlib.pyplot as plt
7
8 # 1. Data Loading
9 transform = transforms.Compose([transforms.ToTensor(), transforms.Normalize((0.5,),
10                                     (0.5,))])
11 batch_size = 64
12
13 trainset = torchvision.datasets.MNIST(root='./data', train=True, download=True,
14                                     transform=transform)
15 trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size, shuffle=
16                                     True)
17
18 testset = torchvision.datasets.MNIST(root='./data', train=False, download=True,
19                                     transform=transform)
20 testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size, shuffle=
21                                     False)
22
23 # 2. Define LeNet-5 Model
24 class LeNet5(nn.Module):
25     def __init__(self):
26         super(LeNet5, self).__init__()
27         self.conv1 = nn.Conv2d(1, 6, kernel_size=5, stride=1, padding=2) # Output: 6
28         # x28x28
29         self.conv2 = nn.Conv2d(6, 16, kernel_size=5, stride=1) # Output:
30         # 16x12x12
31         self.fc1 = nn.Linear(16 * 5 * 5, 120) # Fully
32         # connected layer
33         self.fc2 = nn.Linear(120, 84) # Fully
34         # connected layer
35         self.fc3 = nn.Linear(84, 10) # Output
36         # layer for 10 classes
37
38     def forward(self, x):
39         x = torch.tanh(self.conv1(x)) # First convolution layer
40         # + activation
41         x = torch.nn.functional.avg_pool2d(x, 2) # Average pooling layer (
42         # downsample to 14x14)
43         x = torch.tanh(self.conv2(x)) # Second convolution
44         # layer + activation
45         x = torch.nn.functional.avg_pool2d(x, 2) # Average pooling layer (
46         # downsample to 5x5)
47         x = x.view(x.size(0), -1) # Flatten dynamically
48         x = torch.tanh(self.fc1(x)) # Fully connected layer +
49         # activation
50         x = torch.tanh(self.fc2(x)) # Fully connected layer +
51         # activation
52         x = self.fc3(x) # Output layer
53         return x
54
55 model = LeNet5()
56
57 # 3. Loss Function and Optimizer
58 criterion = nn.CrossEntropyLoss()
59 optimizer = optim.Adam(model.parameters(), lr=0.001)
60
61 # 4. Training the Model
62 num_epochs = 10
63 train_losses = []
64
65 for epoch in range(num_epochs):
66     running_loss = 0.0
67     for images, labels in trainloader:
68         optimizer.zero_grad()
69         outputs = model(images)
70         loss = criterion(outputs, labels)

```

```

55         loss.backward()
56         optimizer.step()
57         running_loss += loss.item()
58
59         average_loss = running_loss / len(trainloader)
60         train_losses.append(average_loss)
61         print(f'Epoch [{epoch+1}/{num_epochs}], Loss: {average_loss:.4f}')
62
63 # 5. Plot Training Loss
64 plt.plot(train_losses, label="Training Loss")
65 plt.xlabel('Epoch')
66 plt.ylabel('Loss')
67 plt.title('Training Loss Over Epochs')
68 plt.legend()
69 plt.show()
70
71 # 6. Evaluate Training Accuracy
72 model.eval()
73 correct_train = 0
74 total_train = 0
75 with torch.no_grad():
76     for images, labels in trainloader:
77         outputs = model(images)
78         _, predicted = torch.max(outputs, 1)
79         total_train += labels.size(0)
80         correct_train += (predicted == labels).sum().item()
81
82 train_accuracy = 100 * correct_train / total_train
83 print(f'Training Accuracy: {train_accuracy:.2f}%)
84
85 # 7. Evaluate Test Accuracy
86 correct_test = 0
87 total_test = 0
88 with torch.no_grad():
89     for images, labels in testloader:
90         outputs = model(images)
91         _, predicted = torch.max(outputs, 1)
92         total_test += labels.size(0)
93         correct_test += (predicted == labels).sum().item()
94
95 test_accuracy = 100 * correct_test / total_test
96 print(f'Test Accuracy: {test_accuracy:.2f}%)

```

### 3 Classify hymenoptera dataset using ResNet18 network trained on ImageNet1K

#### Load the data

```

1 data_transforms = {
2     'train': transforms.Compose([
3         transforms.RandomResizedCrop(224),
4         transforms.RandomHorizontalFlip(),
5         transforms.ToTensor(),
6         transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
7     ]),
8     'val': transforms.Compose([
9         transforms.Resize(256),
10        transforms.CenterCrop(224),
11        transforms.ToTensor(),
12        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
13    ]),
14 }
15

```

```

16 data_dir = 'data/hymenoptera_data'
17 image_datasets = {x: datasets.ImageFolder(os.path.join(data_dir, x), data_transforms[x
    ])}
18         for x in ['train', 'val']}
19 dataloaders = {x: torch.utils.data.DataLoader(image_datasets[x], batch_size=4, shuffle
    =True, num_workers=4)}
20         for x in ['train', 'val']}
21 dataset_sizes = {x: len(image_datasets[x]) for x in ['train', 'val']}
22 class_names = image_datasets['train'].classes
23
24 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")

```

### Visualize data

```

1 def imshow(inp, title=None):
2     """Display image for Tensor."""
3     inp = inp.numpy().transpose((1, 2, 0))
4     mean = np.array([0.485, 0.456, 0.406])
5     std = np.array([0.229, 0.224, 0.225])
6     inp = std * inp + mean
7     inp = np.clip(inp, 0, 1)
8     plt.imshow(inp)
9     if title is not None:
10         plt.title(title)
11     plt.pause(0.001) # pause a bit so that plots are updated
12
13
14 # Get a batch of training data
15 inputs, classes = next(iter(dataloaders['train']))
16
17 # Make a grid from batch
18 out = torchvision.utils.make_grid(inputs)
19
20 imshow(out, title=[class_names[x] for x in classes])

```

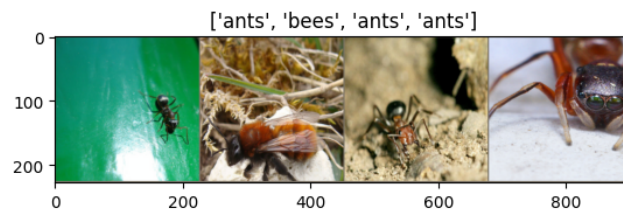


Figure 5: Visualize the data

### Training the model

```

1 def train_model(model, criterion, optimizer, scheduler, num_epochs=25):
2     since = time.time()
3
4     # Create a temporary directory to save training checkpoints
5     with TemporaryDirectory() as tempdir:
6         best_model_params_path = os.path.join(tempdir, 'best_model_params.pt')
7
8         torch.save(model.state_dict(), best_model_params_path)
9         best_acc = 0.0
10
11     for epoch in range(num_epochs):
12         print(f'Epoch_{epoch}/{num_epochs-1}')
13         print('-' * 10)
14
15         # Each epoch has a training and validation phase
16         for phase in ['train', 'val']:
17             if phase == 'train':

```

```

18         model.train() # Set model to training mode
19     else:
20         model.eval() # Set model to evaluate mode
21
22     running_loss = 0.0
23     running_corrects = 0
24
25     # Iterate over data.
26     for inputs, labels in dataloaders[phase]:
27         inputs = inputs.to(device)
28         labels = labels.to(device)
29
30         # zero the parameter gradients
31         optimizer.zero_grad()
32
33         # forward
34         # track history if only in train
35         with torch.set_grad_enabled(phase == 'train'):
36             outputs = model(inputs)
37             _, preds = torch.max(outputs, 1)
38             loss = criterion(outputs, labels)
39
40         # backward + optimize only if in training phase
41         if phase == 'train':
42             loss.backward()
43             optimizer.step()
44
45         # statistics
46         running_loss += loss.item() * inputs.size(0)
47         running_corrects += torch.sum(preds == labels.data)
48     if phase == 'train':
49         scheduler.step()
50
51     epoch_loss = running_loss / dataset_sizes[phase]
52     epoch_acc = running_corrects.double() / dataset_sizes[phase]
53
54     print(f'{phase} Loss: {epoch_loss:.4f} Acc: {epoch_acc:.4f}')
55
56     # deep copy the model
57     if phase == 'val' and epoch_acc > best_acc:
58         best_acc = epoch_acc
59         torch.save(model.state_dict(), best_model_params_path)
60
61     print()
62
63     time_elapsed = time.time() - since
64     print(f'Training complete in {time_elapsed//60:.0f}m {time_elapsed%60:.0f}
65         s')
66     print(f'Best val Acc: {best_acc:4f}')
67
68     # load best model weights
69     model.load_state_dict(torch.load(best_model_params_path, weights_only=True))
70     return model

```

### Visualizing the model predictions

```

1 def visualize_model(model, num_images=6):
2     was_training = model.training
3     model.eval()
4     images_so_far = 0
5     fig = plt.figure()
6
7     with torch.no_grad():
8         for i, (inputs, labels) in enumerate(dataloaders['val']):
9             inputs = inputs.to(device)

```

```

10     labels = labels.to(device)
11
12     outputs = model(inputs)
13     _, preds = torch.max(outputs, 1)
14
15     for j in range(inputs.size()[0]):
16         images_so_far += 1
17         ax = plt.subplot(num_images//2, 2, images_so_far)
18         ax.axis('off')
19         ax.set_title(f'predicted: {class_names[preds[j]]}')
20         imshow(inputs.cpu().data[j])
21
22     if images_so_far == num_images:
23         model.train(mode=was_training)
24         return
25     model.train(mode=was_training)

```

### 3.1 Finetuning the model

```

1 model_ft = models.resnet18(weights='IMAGENET1K_V1')
2 num_ftrs = model_ft.fc.in_features
3 # Here the size of each output sample is set to 2.
4 # Alternatively, it can be generalized to 'nn.Linear(num_ftrs, len(class_names))'.
5 model_ft.fc = nn.Linear(num_ftrs, 2)
6
7 model_ft = model_ft.to(device)
8
9 criterion = nn.CrossEntropyLoss()
10
11 # Observe that all parameters are being optimized
12 optimizer_ft = optim.SGD(model_ft.parameters(), lr=0.001, momentum=0.9)
13
14 # Decay LR by a factor of 0.1 every 7 epochs
15 exp_lr_scheduler = lr_scheduler.StepLR(optimizer_ft, step_size=7, gamma=0.1)

```

#### train and evaluate

```

1 model_ft = train_model(model_ft, criterion, optimizer_ft, exp_lr_scheduler,
2                         num_epochs=25)

```

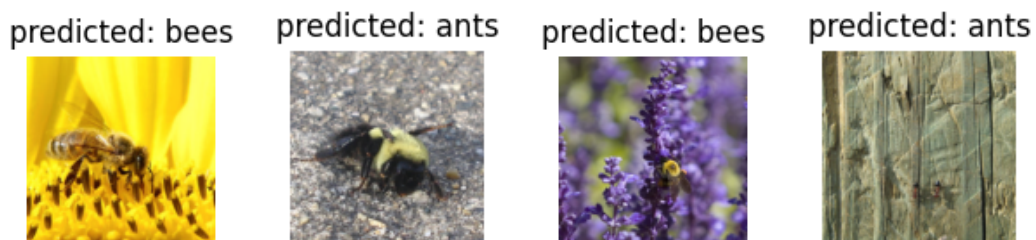


Figure 6: Predicted Images by the model

### 3.2 Use the network as a feature extractor

```

1 model_conv = torchvision.models.resnet18(weights='IMAGENET1K_V1')
2 for param in model_conv.parameters():
3     param.requires_grad = False
4
5 # Parameters of newly constructed modules have requires_grad=True by default
6 num_ftrs = model_conv.fc.in_features
7 model_conv.fc = nn.Linear(num_ftrs, 2)

```



```

8
9 model_conv = model_conv.to(device)
10
11 criterion = nn.CrossEntropyLoss()
12
13 # Observe that only parameters of final layer are being optimized as
14 # opposed to before.
15 optimizer_conv = optim.SGD(model_conv.fc.parameters(), lr=0.001, momentum=0.9)
16
17 # Decay LR by a factor of 0.1 every 7 epochs
18 exp_lr_scheduler = lr_scheduler.StepLR(optimizer_conv, step_size=7, gamma=0.1)

```

### Train and Evaluate

```

1 model_conv = train_model(model_conv, criterion, optimizer_conv,
2                           exp_lr_scheduler, num_epochs=25)

```

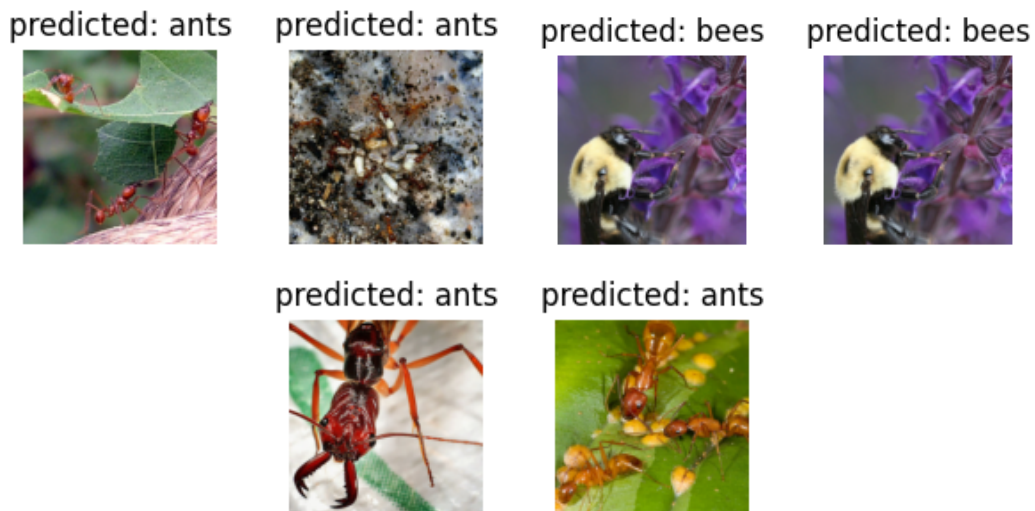


Figure 7: Visualize predicted images using feature extrction

In this task, we applied two common transfer learning techniques using a pre-trained ResNet18 model to classify the Hymenoptera dataset, which consists of images of ants and bees. For fine-tuning, we unfroze the weights of the final layers of the pre-trained model and trained them on our dataset, allowing the model to adjust its parameters specifically for the new task while retaining the learned features from ImageNet. In feature extraction, we froze the weights of the pre-trained model and only trained a new classifier layer on top, using the features extracted by the ResNet18 backbone. These approaches allow the model to leverage the general visual features learned from ImageNet and adapt to the specific task of classifying hymenoptera species with minimal additional training.